

Asymmetries of Vulnerability: How Party Lines Shape Racialized Online Violence Against Women in the 2024 U.S. Election

Jungmin Han¹

Müge Finkel²

Dhanaraj Thakur³

Steven Finkel⁴

Firat Duruşan⁵

Erdem Yörük⁶

Word count: 8,426
(counted w/o abstract and bibliography)

Abstract

Social media platforms have become central to political communication, but they have also intensified online violence against women in politics. Existing research shows that this violence is both gendered and racialized, with women of color facing distinctive forms of abuse. This study extends such an intersectional framework by introducing partisan identity as a key dimension shaping vulnerability to online violence. We argue that party affiliation conditions how race and gender are interpreted in polarized U.S. politics, making Democratic women of color and white Republican women especially visible targets for online hostility. To test this argument, we analyze more than 700,000 tweets mentioning candidates in the 2024 U.S. congressional elections and use transformer-based classifiers to detect multiple types of offensive and hate speech. The findings reveal asymmetric patterns of online violence. Democratic women of color, especially Black and Asian women, face elevated levels of abuse compared with both their male co-partisans and white Democratic women, while white Republican women experience heightened exposure compared with both their male co-partisans and Republican women of color. These patterns are evident across multiple forms of online violence, from broader offensive speech to more specific identity-targeted hate speech. By showing how party affiliation intersects with race and gender, this study advances research on online political violence, intersectionality, and democratic representation.

¹ Postdoctoral Researcher, Department of Political Science, Trinity College Dublin

² Associate Professor, School of Public and International Affairs, University of Pittsburgh

³ Director, Emerging Technology Initiative-Multiracial Democracy Project, George Washington University

⁴ Daniel H. Wallace Professor, Department of Political Science, University of Pittsburgh

⁵ Research Assistant, Center for Computational Social Sciences, Koç University

⁶ Professor, Center for Computational Social Sciences, Koç University

Introduction

The widespread use of interactive technologies and social media platforms has transformed the political landscape by enabling low-cost, real-time communication beyond traditional political networks. Initially, these platforms were expected to democratize public discourse and expand political participation, especially for groups historically disadvantaged by party and campaign structures, including women (Dahlgren, 2005). Yet two decades of research across different national contexts has cast doubt on these early expectations, showing that digital platforms have also become spaces where abuse, harassment, hate speech, and misinformation thrive (Amnesty International, 2018; Ward, 2016; Bardall et al., 2020). For women in politics, this online violence poses serious psychological, physical, and democratic harms, contributing to self-censorship, withdrawal from public debate, and even exit from electoral competition (Krook, 2020; Wilfore, 2022; Carlson, 2019; Sobieraj, 2020). Importantly, these harms are not evenly distributed among women. Politically active women face “many axes of discrimination,” including race, religion, sexuality, and political ideology, making some groups especially vulnerable to digital abuse (Kuperberg, 2018; Bardall et al., 2020). Women of color, in particular, experience distinctive forms of online violence shaped by the intersection of racism and misogyny, including misogynoir directed at Black women (Bailey, 2021; Kwarteng et al., 2023). Empirical studies similarly show that women of color candidates, especially Black women, are more likely than white women and men of color to be targeted by misinformation, disinformation, and violent rhetoric online (Norwood et al., 2021; Thakur et al., 2022).

Yet, while this literature has substantially advanced our understanding of the intersectional nature of online political violence, it has paid less attention to how these vulnerabilities are shaped by partisan context. In contemporary American politics, race and gender are not experienced,

interpreted, or mobilized outside of party politics. Partisan polarization increasingly structures how racial and gender identities are represented and contested in the public sphere, with the Democratic and Republican parties attaching different symbolic meanings to diversity, feminism, racial inclusion, and traditional gender hierarchies (Cassese and Barnes, 2019; Junn and Masuoka, 2020; Montoya et al., 2022; Schaffner, 2022; Kim and Junn, 2024). This is especially important because women of color are not distributed evenly across the two major parties, nor are they interpreted in the same way when they run as Democrat or Republican candidates. As a result, party affiliation may intensify, redirect, or moderate the online violence directed at women candidates. Despite this, partisanship has rarely been examined as an intersecting political identity that conditions how race and gender shape exposure to online abuse.

To address this gap, we investigate how race and partisanship intersect to shape online violence against women in contemporary American politics. Our argument is not simply that women candidates face more online violence than men, but that the magnitude of this gendered vulnerability varies across racial and partisan contexts. Among white candidates, we expect the positive effect of being a woman on exposure to online violence to be stronger for Republicans than for Democrats, because white Republican women are often associated with conservative gender politics, traditional femininity, and opposition to feminist or egalitarian claims. Among non-white candidates, by contrast, we expect this gendered effect to be stronger for Democrats than for Republicans, because Democratic women of color are especially visible symbols of racial inclusion, women's political empowerment, and progressive partisan change. Put differently, within the Democratic Party, the gendered vulnerability to online violence should be more pronounced among non-white candidates than among white candidates, whereas within the Republican Party, it should be more pronounced among white candidates than among non-white

candidates. We therefore expect online violence against women candidates to follow an asymmetrical pattern, with partisan affiliation shaping which racialized groups of women become most vulnerable relative to their male counterparts.

To test this argument, we analyze tweets mentioning Democratic and Republican candidates in the 2024 U.S. congressional elections. Our research design centers on a measurement strategy that identifies online violence at the tweet level and links it to candidates' gender, race, and party affiliation. We construct a dataset of 954 House and Senate candidates and collect 759,259 tweets mentioning these candidates between May and August in 2024. To measure online violence, we develop a multi-layered ontology that distinguishes offensive speech from identity-targeted hate speech, treating hate speech as a subset of offensive language directed at candidates based on race, ethnicity, religion, or sexuality. Offensive speech was also measured across multiple dimensions, including profanity, demeaning language, and ridicule. We then fine-tune transformer-based classifiers using a human-annotated gold-standard corpus of 3,690 tweets to identify offensive and hateful content systematically across the full dataset. Finally, we estimate logistic regression models predicting whether tweets contain offensive or hate speech as a function of candidate gender, race, party affiliation, and their interactions, while controlling for candidates' institutional position, incumbency, online visibility, campaign activity, tweet impressions, and proximity to Election Day. This design allows us to examine whether online violence against women candidates reflects a uniform gendered pattern or varies across racial and partisan contexts.

The findings provide broad support for our hypotheses by showing that the gendered effect of online violence varies across racial and partisan contexts. Among white candidates, women's heightened exposure to online violence is concentrated among Republicans: white Republican women face an approximately 4.2 percentage-point higher predicted probability of offensive

speech than white Republican men, while the corresponding effect among white Democrats is negative. Among non-white candidates, the pattern is reversed, with non-white Democratic women experiencing about a 2.5 percentage-point higher predicted probability of offensive speech than non-white Democratic men, whereas the estimates for non-white Republicans are smaller, negative, or statistically uncertain. Further disaggregation shows that these vulnerabilities are especially pronounced among Asian and Black Democratic women, who have among the highest predicted probabilities of offensive and hate speech, while white Republican women emerge as the most consistently targeted group within the Republican Party. Together, these results show that online violence against women candidates cannot be explained by gender or race alone. Rather, partisan identity conditions how gender and race are interpreted, producing asymmetric vulnerabilities for Democratic women of color and white Republican women.

This paper makes both theoretical and methodological contributions to the study of online violence against women in politics. Theoretically, it extends scholarship on online political violence as a gendered and intersectional phenomenon (Kuperberg, 2018; Bardall et al., 2020; Krook, 2020; Thakur et al., 2022) by conceptualizing party affiliation as an intersecting political identity that conditions candidates' vulnerability to abuse. Prior research shows that women of color, especially Black women, face distinctive forms of digital hostility rooted in racism and misogyny (Bailey, 2021; Norwood et al., 2021; Kwarteng et al., 2023), but less attention has been paid to how these vulnerabilities vary across partisan contexts. Rather than treating women, women of color, or political candidates as homogeneous groups, this paper argues that gendered vulnerability is shaped by candidates' race-party location within polarized U.S. politics. Methodologically, we build on existing approaches to measuring online abuse in politics (Rheault et al. 2019; Erikson et al. 2023; Holm et al. 2024) by moving beyond aggregate or binary measures

of abusive content to differentiate among multiple forms of offensive and hate speech. Using supervised machine-learning techniques, we show that online violence appears in diverse forms, ranging from ridicule and demeaning language to identity-targeted attacks based on gender, race, and religion.

Gendered and Racialized Online Violence against Women in Politics

Despite “the obvious positive consequences that the Internet extends and pluralizes the public sphere in a number of ways - [including its alignment with] the idea that democracy resides ultimately, with citizens who engage in talk with each other,” the promise of social media platforms for democratization and equalization purposes have been short lived (Dahlgren, 2005). Instead, these platforms have evolved into spaces where violent, harassing, and discriminatory language thrives, disproportionately targeting women and girls — with 1 in 10 women over the age of 15 having experienced some form of online violence. (European Institute for Gender Equality, 2017). Furthermore, women who are actively involved in politics emerge as the second largest target group of online violence after women journalists — 27 times more likely to face abuse online than their male counterparts (Amnesty International, 2018).

As women’s participation and representation in politics increase, unfortunately, so does their risk of becoming victims of online violence. Women breaking the glass ceiling by achieving high levels of public recognition in politics are more likely to receive uncivil messages than their male counterparts (Rheault et al., 2019; Han et al., 2026). The consequences of this growing and pervasive threat to women’s equal rights to political and public participation globally are immense. At the macro level, in addition to the dangers online violence poses to gender equality, it also

threatens the legitimacy of political and electoral processes and poses unprecedented challenges to the quality of democratic governance (Haciyakupoglu & Wong, 2024; Phungula, 2023; Wilfore, 2022). The U.N. Human Rights Council finds the erosion of women’s human rights to be “a litmus test for the human rights standards of the society as a whole,” points to the broader ramifications of the tech-enabled backlash against women’s rights for global peace and security (UN Human Rights Council, 2018). At the individual level, online violence against women in politics presents a distinct form of gender-based violence by targeting women because of their gender, and by taking on gendered forms (Krook, 2020). As such, online violence perpetuates and expands the types of gender-based violence women politicians already face, including physical, sexual, and psychological, through new mediums and tools, such as doxing, misinformation and disinformation (Krook, 2020; Bardall, 2023; Wilfore, 2022). Finally, it further disenfranchises both incumbent politicians and aspirants, known as the ‘silencing effect’ discourages women from being and becoming politically active (Amnesty International, 2018; Bardall, 2023; Carlson, 2019).

Existing research shows that online violence is not only distinctly gendered, but also deeply intersectional. Rather than targeting women as a uniform group, online abuse often attaches to women’s layered and multiple identities, placing some women at greater risk than others. Politically active women, in particular, face “many axes of discrimination” on the basis of ethnicity, religion, ideology, race, sexuality, disability, and other identity markers (Kuperberg, 2018; Bardall et al., 2020). These forms of intersectional discrimination reproduce broader societal stereotypes (Zetterberg et al., 2022) and can contribute to women’s continued underrepresentation in political decision-making positions (Zamfir, 2024). In the U.S. context, studies also document rising levels of incivility across online platforms and political discourse, with female and minority candidates

disproportionately affected by online hostility and abuse (Thakur et al., 2022; Rheault et al., 2019; Kuperberg, 2022).

An intersectional approach is especially useful for understanding how different groups engage with, and are treated within, political systems. It allows scholars to examine how individual identities shape political participation, representation, and political vulnerability in ways that cannot be captured by looking at gender, race, or other categories in isolation. Crenshaw's (1991) seminal work established that individuals can occupy multiple marginalized identities, producing forms of oppression that are not merely additive but qualitatively distinct. This insight is particularly important for understanding the experiences of women of color, whose participation in representative politics differs from that of white women, men of color, and white men. Scholars of gender and politics have therefore argued that intersectionality is critical to the study of politics, despite the conceptual and methodological challenges it presents (Smooth, 2006). By drawing attention to differences both between groups, such as men and women, and within groups, such as differently situated women, intersectional analysis moves research away from gender essentialism (Jordan-Zachery, 2007). It also offers methodological advantages over analyses that treat race and gender separately, because it requires scholars to develop theories and approaches that account for how multiple identity categories jointly shape political participation and outcomes (Simien, 2007).

Building on this framework, scholars of representative politics have examined how women of color encounter distinctive forms of exclusion, marginalization, and resistance. Hawkesworth (2003), for example, shows how racism and misogyny have combined to exclude and silence women of color in the U.S. Congress. Similarly, Brown et al. (2023) examine how race, gender, and other intersecting identities shape the experiences of women of color in politics, including women serving in representative roles beyond Congress. This body of work demonstrates that

women of color do not simply experience the barriers faced by women and racial minorities separately; rather, they confront forms of political marginalization produced through the interaction of race, gender, and other identities.

More recent scholarship extends these concerns to digital spaces, where social media has become central to political participation, campaigning, and public visibility. In this context, the concept of “misogynoir,” coined by Moya Bailey (2021), captures a specific form of hostility directed at Black women through the intertwined logics of racism and sexism. Because Black women’s lived experiences shape their political perspectives and contributions to public debate, misogynoir also highlights the particular challenges they face when using digital platforms to participate in politics and advocate for social change (Kwarteng et al., 2023).

Empirical studies of women of color in digital politics further show how these intersectional vulnerabilities operate in practice. Norwood et al. (2021), for instance, interviewed Black women politicians to examine their experiences online and the relationship between digital hostility and the underrepresentation of Black women in political office. Thakur et al. (2022) similarly conducted a content analysis of Twitter, now X, posts and found that women of color—and especially Black women—running for Congress in 2020 were twice as likely to be targeted by mis- and disinformation and four times more likely to experience violent abuse compared with white women and men of color. At the same time, their work shows that women of color developed strategies of resilience to continue campaigning despite these attacks. Together, this literature suggests that online violence against women in politics must be understood through an intersectional lens, with particular attention to how race and gender jointly structure political visibility, vulnerability, and resistance in digital spaces.

Although existing research has made important contributions to understanding the gendered and intersectional nature of online violence against women in politics, less is known about how party affiliation conditions these dynamics. In particular, there is still no systematic account of whether political party, as a form of political identity, intersects with race and gender to shape candidates' vulnerability to online abuse. This gap is important because party affiliation is not merely an institutional label or campaign characteristic. In the contemporary United States, party functions as a socially meaningful identity that organizes perceptions of race, gender, ideology, and political belonging. We therefore build on intersectionality scholarship to examine how the intersecting identities of political candidates are associated with online violence, while placing particular emphasis on the moderating role of party affiliation.

How Party Lines Shape Racialized Online Violence against Female Electoral Candidates

We argue that partisanship is critical for understanding why women candidates may be targeted in different ways depending on their racial identity. In American politics, gender and race are not only salient identities — they are increasingly structured by partisan polarization. The Democratic and Republican parties do not treat these identities symmetrically. Democratic elites and voters have, over time, adopted a platform of multiculturalism, diversity, and inclusion, embracing women of color as symbols of progress and as representatives of the party's commitment to historically marginalized groups. This orientation is reflected in both rhetoric and recruitment patterns, particularly since the 2018 and 2020 election cycles (Montoya et al. 2022; Schaffner 2022). Meanwhile, Republican discourse has increasingly leaned into appeals that resonate with ethno-

nationalist sentiment, with Trump’s campaigns in 2016 and 2020 amplifying racial resentment and hostile sexism as mobilizing tools (Bracic et al. 2019; Cassese and Barnes 2019; Schaffner 2022; Kim and Junn 2024). These divergent ideological positions place women of color at the intersection of partisan identity formation and racialized gender discourse.

Additionally, the notion of a cohesive “women’s vote” has long obscured stark racial divisions within the electorate. While women overall are more likely than men to vote for Democratic candidates, this gender gap is, in large part, a reflection of the overwhelming support that women of color—particularly Black women—consistently give to the Democratic Party (Junn and Masuoka 2020; Bejarano 2013; Harris-Perry 2011). In contrast, white women have voted in the majority for Republican presidential candidates in nearly every election since the 1950s, with the only exceptions being 1964 and 1996 (Junn 2017). Research shows that support for Donald Trump among white women in 2016 and 2020 was strongly predicted by the endorsement of sexist beliefs and diminished perceptions of discrimination against women (Bracic et al. 2019; Cassese and Barnes 2019). These patterns suggest that gender alone cannot explain vote choice; rather, the intersection of race and gender—mediated by ideology—plays a pivotal role in shaping political behavior.

Intersecting identities provides a critical lens for understanding how political identities are constructed and mobilized in partisan contexts. As Montoya et al. (2022) and Kim and Junn (2024) demonstrate using multiracial post-election surveys, shared racial and gender identities significantly shape perceptions of representation, but their effects vary depending on the group in question and the broader political environment. For African American and Latina women, descriptive representation within the Democratic Party signals greater inclusion and empowerment, a pattern that aligns with the “minority empowerment” thesis (Bobo and Gilliam 1990; Pantoja

and Segura 2003). At the same time, the absence or delegitimization of women of color in Republican politics contributes to political alienation and signals exclusion. Voters often assume that women and minority candidates are Democrats (Gershon and Monforti 2021), reinforcing partisan stereotypes and making women of color, especially those in Democratic roles, potent symbols of ideological and demographic change.

These dynamics also help explain why women of color aligned with the Democratic Party attract outsized attention—both positive and negative—in the public sphere. Their symbolic role as challengers to traditional hierarchies of race and gender makes them especially salient targets for backlash, particularly in a political climate increasingly shaped by the fear of demographic replacement and loss of status among segments of the white electorate (Junn 2017; Junn and Masuoka 2020). As such, Democratic women of color are not simply treated as political actors but as existential threats to a social order rooted in white patriarchal dominance. This perception, deeply tied to the preservation of “white womanhood” as a political identity (Junn 2017; Collins 1990), enables the justification of racialized and gendered attacks—rhetorical or otherwise—against these figures. This pattern also resonates with findings that suggest that women of color politicians who explicitly advocate feminist strategies are further targeted with online violence rhetoric (Banet-Weiser and Miltner 2016, Lewis et al., 2017).

In line of this view, while women of color across the political spectrum are often subject to gendered and racialized marginalization, the intensity and nature of their exposure to online hate speech are deeply contingent on their partisan affiliation. Women of color running as Democratic candidates are likely to be significantly more vulnerable to online political violence than their white Democratic counterparts. While white Democratic women may also face online misogyny—particularly when advancing feminist rhetoric—they are often spared the racialized abuse that

uniquely targets women of color. Moreover, white Democratic women can more easily be integrated into mainstream political narratives without triggering the same intersectional backlash. The combination of being a woman, a person of color, and a progressive voice places Democratic women of color at the epicenter of racialized and gendered vitriol in digital spaces.

On the other hand, within the Republican Party, it is white women candidates who are likely to face more online political violence than their women-of-color counterparts. Republican white women often embody traditional, submissive femininity or explicitly support policies perceived as hostile to gender equality—stances that can provoke intense criticism, particularly from liberal and progressive voters. Their visible alignment with patriarchal authority and anti-feminist rhetoric renders them highly legible targets for those who oppose Republican ideology and its gender politics. In contrast, Republican women of color may be more shielded from online attacks, not because they are spared misogyny or racism, but because their racial identity intersects with broader Democratic values of diversity and multicultural inclusion. For Democratic partisans and progressive critics, the racial identity of Republican women of color can produce ambivalence—limiting the extent of vilification they receive or at least moderating the tone of attacks. Their identity may even complicate simplified narratives of Republican regression, making them less prominent figures in online hate campaigns. In this way, race can serve as a moderating force, reducing the visibility and intensity of political violence against Republican women of color compared to their white counterparts.

Taken together, these arguments suggest that online violence against women candidates should not be understood as a uniform gendered phenomenon. Rather, the effect of being a woman on exposure to online violence is likely to vary depending on how gender is situated within racial and partisan contexts. Party affiliation changes the symbolic meaning of race and gender: the same

racialized gender identity may be read differently depending on whether a candidate runs as a Democrat or a Republican. Likewise, the same partisan affiliation may generate different vulnerabilities depending on whether women candidates are white or non-white. We therefore conceptualize candidates' vulnerability through a two-by-two framework that combines race and party, as shown in Table 1.

Table 1. Expected Patterns of Online Gendered Violence by Race and Party

	Republican candidates	Democratic candidates
White candidates	White Republican women are expected to face heightened gendered hostility because they are often associated with conservative gender politics, traditional femininity, and opposition to feminist or egalitarian claims.	White Democratic women may face gendered hostility, especially when associated with feminist or progressive claims, but they are less likely to be targeted through racialized abuse.
Non-white candidates	Republican women of color may face racism and misogyny, but their racial identity may also complicate attacks from progressive critics and reduce their visibility as targets.	Democratic women of color are expected to face heightened gendered hostility because they symbolize racial inclusion, women's empowerment, and progressive partisan change.

This framework implies two sets of comparisons. First, within each racial group, party affiliation should condition the gendered effect of online violence. Among white candidates, the woman-versus-man difference in exposure to online violence should be larger for Republicans than for Democrats, because white Republican women are more likely to be targeted for their association with conservative gender politics. Among non-white candidates, the pattern should reverse: the woman-versus-man difference should be larger for Democrats than for Republicans,

because Democratic women of color are especially visible symbols of racial and gender progress within contemporary partisan conflict. Second, within each party, race should condition the gendered effect of online violence. Among Democratic candidates, the woman-versus-man difference should be larger for non-white candidates than for white candidates, because Democratic women of color are exposed to both misogynistic and racialized backlash. Among Republican candidates, however, the woman-versus-man difference should be larger for white candidates than for non-white candidates, because white Republican women are more legible targets for criticism of conservative gender ideology, whereas Republican women of color may occupy a more ambivalent position in partisan discourse.

These arguments lead to the following testable hypotheses:

- Hypothesis 1:** Among white candidates, the positive effect of being a woman on exposure to online violence will be stronger for Republican candidates than for Democratic candidates.
- Hypothesis 2:** Among non-white candidates, the positive effect of being a woman on exposure to online violence will be stronger for Democratic candidates than for Republican candidates.
- Hypothesis 3:** Among Democratic candidates, the positive effect of being a woman on exposure to online violence will be stronger for non-white candidates than for white candidates.
- Hypothesis 4:** Among Republican candidates, the positive effect of being a woman on exposure to online violence will be stronger for white candidates than for non-white candidates.

Research Design

For our empirical analysis, we examine online violence on X, formerly Twitter, directed at candidates in the 2024 U.S. congressional elections.⁷ The 2024 election provides a particularly appropriate case for testing our argument for several reasons. First, it took place in a highly polarized political environment in which race, gender, and party were tightly linked in public discourse. This was especially visible in the simultaneous presidential contest, where the Democratic nominee, Kamala Harris, was a Black and South Asian woman, while the Republican nominee, Donald Trump, was a white man. As a result, the national campaign environment made questions of gender, race, representation, feminism, and partisan identity unusually salient. These dynamics were further shaped by the Supreme Court’s decision in *Dobbs v. Jackson Women’s Health Organization*, in which abortion rights and gender equality remained central campaign issues (Evans, 2026). Second, the 2024 cycle revealed clear partisan asymmetries in women’s congressional candidacies. CAWP (2024) reported that the number of women congressional nominees fell short of recent cycles and that this decline was concentrated entirely in the Republican Party, while Democratic women performed especially strongly in House and Senate primaries. This makes 2024 a useful context for examining whether the political meaning of women’s candidacy differs across party lines. Third, the election took place in an online environment in which X remained especially central to political communication, with Pew Research Center’s (2024) finding that most X users used the platform to follow politics or political issues.

⁷ Our dataset was first introduced at the Companion Proceedings of The ACM Web Conference 2025 (Finkel et al., 2025).

We began our data collection by compiling a comprehensive list of 954 Congressional candidates from both the Democratic and Republican Parties—covering races for both the House and Senate—drawing on publicly available sources such as Ballotpedia. For each candidate, we recorded their name, gender, race, party affiliation, state, incumbency status, follower count, tweet count, and X’s username. Notably, Table 2 shows clear differences between Democratic and Republican congressional candidates in terms of both gender and race. Although the total number of Democratic and Republican candidates is relatively similar, with 471 Democrats and 483 Republicans, their demographic composition differs substantially. Republican candidates are overwhelmingly male: 393 of 483 Republican candidates, or approximately 81%, are men, while only 90, or about 19%, are women. By contrast, Democratic candidates are more gender-balanced, with women comprising 206 of 471 candidates, or about 44%. The racial distribution also differs by party. Republican candidates are predominantly white, with 394 of 483 candidates, or about 82%, identified as white. Democratic candidates are more racially diverse, with non-white candidates making up 185 of 471 candidates, or about 39%, compared with only 18% among Republicans.

Table 2. Distribution of 2024 U.S. Congress Candidates by Gender, Race, and Party

	Male (n = 658)		Female (n = 296)		Total
	Republican	Democrat	Republican	Democrat	
White	330	163	64	123	680
Non-White	63	102	26	83	274
Total	393	265	90	206	954

Next, we gathered all tweets posted between May 20 and August 23, 2024,⁸ which mentioned any of these candidates with at least one active Twitter/X account. Tweets referencing the candidates were gathered at a time when academic access to the platform’s official API had become highly restricted. To address these constraints, recent computational studies have relied on the Nitter web interface (Alzahrani & AlGhamdi 2025; Ortiz et al. 2025; Han et al. 2026). Nitter is an open-access, privacy-oriented alternative front end that allows researchers to access and extract publicly available content from X without requiring authenticated API access (www.nitter.net). Employing Nitter facilitates the systematic and reproducible collection of public political communication while minimizing exposure to platform-imposed access limitations and algorithmic personalization effects. To improve data quality, we excluded tweets that were not in English or that included more than three user mentions, as these were more likely to be spam or unrelated content.

The final dataset comprised 759,259 tweets.⁹ Table 3 presents the distribution of tweets mentioning candidates by gender, race, and party, with the average number of tweets per candidate reported in parentheses. Overall, tweets mentioning male candidates account for a larger share of the dataset than tweets mentioning female candidates, with 496,413 tweets mentioning men compared to 244,359 tweets mentioning women. However, the average number of tweets per candidate (in parenthesis) reveals important differences across race–party groups. Among white candidates, Republican women receive the highest average number of tweets per candidate, with 1,130 tweets, compared to 799 tweets for white Republican men, 732 tweets for white Democratic

⁸ This window captures a period when congressional candidates were publicly visible and when partisan, racial, and gendered identities were already salient in the campaign environment. Because the data do not include September through Election Day, our findings should be interpreted as evidence of online violence during the pre-general-election campaign period rather than as a complete account of the entire 2024 election cycle.

⁹ The discrepancy between this count and the overall total in Table 3 reflects candidates who were affiliated with neither the Republican nor the Democratic Party ($n = 18,487$).

men, and 674 tweets for white Democratic women. Among non-white candidates, Democratic women receive the highest average number of tweets per candidate, with 1,039 tweets, exceeding non-white Democratic men, who receive 731 tweets on average. By contrast, non-white Republican women receive far fewer tweets in total and on average, with 2,872 total tweets and an average of 836 tweets per candidate. These patterns suggest that public attention on Twitter/X is not evenly distributed across candidates. Instead, it varies substantially by the intersection of gender, race, and party, with white Republican women and non-white Democratic women receiving especially high average levels of attention.

Table 3. Distribution of Tweets Targeting Candidates by Gender, Race, and Party

	Male (n = 496,413)		Female (n = 244,359)		Total
	Republican	Democrat	Republican	Democrat	
White	263,780 (799)	119,247 (732)	72,321 (1,130)	82,952 (674)	538,300
Non-White	38,780 (616)	74,606 (731)	2,872 (836)	86,214 (1,039)	202,472
Total	302,560	193,853	75,193	169,166	740,772

Note: The first number in each cell represents the total count of tweets mentioning the candidate categorized within the group. The number in parentheses indicates the average number of tweets mentioning an individual candidate from that group.

The core analytical task of this study is to identify whether tweets referencing electoral candidates contain offensive or hateful language. Offensive and hateful speech constitute one form of online violence against women in politics, alongside cyberstalking, sexual harassment, doxing, intimate image abuse, and disinformation (Krook, 2020). Such language can attack, intimidate, or silence women candidates directly, while sexist and misogynistic tropes may also operate as implicit “dog whistles” that enable harassment while preserving plausible deniability. Because

these forms of abuse can be both explicit and context-dependent, their detection requires categories that are conceptually flexible yet clearly delimited. We therefore construct a multi-layered ontology of offensive and hateful speech that captures distinct forms of violence while remaining applicable as document-level labels.

We define hate speech using the framework provided by the United Nations Strategy and Plan of Action on Hate Speech, which describes it as any form of spoken, written, or behavioral communication that attacks individuals or groups, or relies on pejorative or discriminatory language, because of their identity. These identities may include religion, ethnicity, nationality, race, color, descent, gender, or other identity-based characteristics (United Nations, 2020, pp. 7–8). This definition informs our classification framework by emphasizing two analytically distinct components: the form of harmful expression, such as insult, ridicule, demeaning language, or profanity, and the identity category being targeted, such as gender, ethnicity, race, religion, or sexuality. Rather than relying on a single general model to detect all forms of abuse, we develop separate classification tools for different components of offensive and hateful language and then combine these components into broader measures.

Our ontology distinguishes between offensive speech and hate speech. *Offensive speech* refers to words or phrases that assault, threaten, insult, ridicule, or belittle a specific individual, group, or institution. Within this broader category, we identify three subtypes. *Ridicule* refers to language that mocks, shames, or ridicules an individual or group; although such language may appear harmless in familiar settings, in public contexts it can function as bullying and cause harm. *Demeaning* speech refers to language that degrades or belittles a target’s character, moral worth, or humanity, including discriminatory language. *Profanity* refers to insulting language that uses obscenities or expressions considered morally repugnant.

Hate speech is treated as a subset of offensive speech that targets an individual or social group on the basis of identity, including gender, ethnicity, race, religion, or sexuality. We distinguish among three types of hate speech. *Religious hate* refers to offensive speech directed at individuals or groups based on religious belief, non-belief, alleged heresy, or membership in a faith community or denomination. *Ethnic hate* refers to offensive speech targeting ethnic, cultural, racial, or national identity. *Sexuality hate* refers to offensive speech targeting gender identity, sexual identity, or gendered social position, including sexist language, discriminatory language, and gender stereotyping.

Building our automated detection strategy around these related concepts offers several advantages. First, the ontology allows us to distinguish different levels and forms of severity, ranging from direct threats and insults to less overt but still harmful forms of ridicule or degradation. Second, it allows us to examine not only whether candidates are exposed to offensive or hateful speech, but also whether different groups of candidates are exposed to different types of online violence. Finally, by detecting identity-targeted speech across gender, religious, ethnic, racial, and sexuality-based categories before combining them into a broader measure of hate speech, the ontology reduces the risk of excluding groups targeted through distinct forms of abuse. In this sense, our framework is flexible enough to capture the multidimensional character of online violence and precise enough to generate analytically meaningful categories for empirical analysis.

While conducting multiple detection tasks of these concepts, we employed human-annotated gold standard data to fine-tune BERT-based transformer models. Fine-tuning with human-verified annotations allows the models to learn the specific linguistic and contextual patterns present in politically charged discourse, improving their accuracy for nuanced

classification. To build our training dataset, we annotated 4,000 tweets with the help of four domain experts. Given the rarity of offensive and hateful content in randomly sampled data, we adopted an active learning approach to enrich the dataset with relevant examples while minimizing systematic recall bias. An initial set of five high-recall models—including ChatGPT-4.0 and traditional classifiers—was trained on a small, random sample of annotated tweets. These models were then used to scan a larger pool of tweets and identify likely instances of hate speech. From this, 3,000 tweets flagged by at least one model were selected for annotation. These were combined with 1,000 randomly sampled tweets, resulting in a cleaned and adjudicated gold-standard corpus of 3,690 tweets, including 959 labeled as offensive and 277 labeled as hate speech.

Using this corpus, we developed a series of binary classification models to identify offensive and hate speech in the dataset. For offensive speech detection, we fine-tuned the `'cardiffnlp/twitter-roberta-base-offensive'` model (Vidgen et al. 2021) using our annotated data, achieving strong performance with an F1 macro score of 0.89 on the validation set and 0.90 on the test set. To classify specific forms of offensive speech, we trained separate models for each subcategory rather than using a single multiclass multilabel model. This approach was intended to improve performance given the limited size of the annotated dataset and to allow each model to better capture the distinct linguistic features associated with each category. Specifically, we fine-tuned `'xlm-roberta-base'` to detect profanity, achieving an F1 macro score of 0.95 on the test set; `'xlm-roberta-base'` to detect demeaning language and ridicule, achieving F1 macro scores of 0.79 and 0.83, respectively. For hate speech detection, we fine-tuned the `'facebook/roberta-hate-speech-dynabench-r4-target'` model as a binary classifier (Camacho-Collados et al., 2022). To identify specific types of hate speech, we used the pretrained `'wesleyacheng/hate-speech-`

multilabel-classification-with-bert` model, focusing on the race, gender, and religion labels most relevant to our research questions.


To examine how gender, party affiliation, and race intersect to shape candidates' exposure to offensive or hate speech, we estimate a logistic regression model where the dependent variable is a binary indicator coded as 1 if a tweet is classified as offensive/hate speech and 0 otherwise. The key independent variables are candidate gender, party affiliation, and race, along with all two-way and three-way interaction terms. This specification allows us to assess whether and how the effect of one identity dimension (e.g., gender) varies depending on the others (e.g., race and party).

Moreover, we include several control variables to account for differences in candidates' institutional position, visibility, and campaign activity that may independently shape exposure to online violence. First, we control for whether a candidate ran for the Senate rather than the House, since Senate campaigns typically receive greater media attention, cover larger constituencies, and may generate higher levels of online discussion. We also include incumbency status, as incumbents generally have greater name recognition, stronger public profiles, and longer-standing partisan reputations than challengers (Gorrell et al. 2020). To account for candidates' online visibility, we control for follower count, since candidates with larger audiences are more likely to attract both supportive and hostile engagement (Rheault et al. 2019; Han et al. 2026). We further control for candidates' own tweet count, which captures campaign activity on X and the extent to which candidates make themselves visible in online discourse. Impression count is included to account for the visibility of the collected tweets mentioning each candidate, as posts that reach larger audiences may be more likely to reflect or amplify public attention, hostility, and potential abuse directed at the candidate. Finally, we include date, measured as proximity to Election Day, because campaign attention and partisan hostility are likely to intensify as the election approaches.

Results

We begin the analysis by estimating the overall association between candidate gender and exposure to online violence. To make the results substantively interpretable, we calculate predicted probabilities from logistic regression models while holding the control variables—Senate, incumbency, follower count logged, tweet count logged, impression count logged, and time—at their mean values. As shown in Figure 1, female candidates are more likely than male candidates to be mentioned in tweets containing both offensive speech and hate speech. The predicted probability of offensive speech is 15.4% for female candidates, compared with 14.5% for male candidates. The gender difference is more pronounced for hate speech. The predicted probability is 1.2% for female candidates, compared with 0.2% for male candidates: female candidates are found to face hate speech at roughly six times the rate of male candidates.

Figure 1. Predicted Probability of Offensive and Hate Speech by Gender




Note: The predicted probabilities of these results are derived from logistic regression models, with other control variables—such as Senate, Incumbency, Follower Count (log), Tweet Count (log), Impression count (log), and Time to Election—held at their mean values.

We next examine whether the gendered pattern of online violence varies across racial groups. As shown in the top panel of Figure 2, the predicted probability of offensive speech is higher for women than men among white, Black, and Asian candidates, but the size of this gender gap differs substantially across groups. The difference is especially pronounced among Asian candidates, where the predicted probability of offensive speech is 21.2% for women compared with 11.5% for men. Black women also face higher predicted levels of offensive speech than Black men, at 17.5% compared with 13.6%. Among white candidates, the gender difference is smaller, with predicted probabilities of 15.1% for women and 14.5% for men. The pattern is not uniform, however: Hispanic women have a lower predicted probability of offensive speech than Hispanic men, 11.8% compared with 18.9%, while candidates in the “Other” racial category show virtually no gender difference, at 17.6% for women and 17.7% for men.

The bottom panel shows that gender differences are more consistently visible for hate speech. Across all racial groups, women candidates have higher predicted probabilities of hate speech than men. The largest gap appears among Black candidates, where the predicted probability is 2.5% for women compared with 0.4% for men. White women also face substantially higher predicted hate speech than white men, at 0.8% compared with 0.1%. Smaller but still positive gender gaps appear among Hispanic candidates, 0.5% for women compared with 0.4% for men; Asian candidates, 1.0% compared with 0.4%; and candidates in the “Other” category, 1.1% compared with 0.5%. These results suggest that gendered online violence is not distributed uniformly across racial groups. While women generally face higher predicted exposure to hate speech, the magnitude of this vulnerability is especially pronounced among Black women and, for offensive speech, among Asian and Black women.

Figure 2. Predicted Probability of Offensive and Hate Speech by Gender and Race



Note: The predicted probabilities of these results are derived from logistic regression models, with other control variables—such as Senate, Incumbency, Follower Count (log), Tweet Count (log), Impression count (log), and Time to Election—held at their mean values.

While Figure 2 shows that the gendered pattern of online violence differs across racial groups, our hypotheses further suggest that these differences should depend on candidates' partisan affiliation. Figure 3 examines the gendered effect of online violence across both candidate race

and party. The estimates show the difference in predicted probability between women and men within each race–party group; positive values indicate that women are more likely than men to be targeted. Overall, the results indicate that the gendered effect of online violence varies substantially across racial and partisan contexts.

First, among white candidates, the gendered effect is considerably stronger for Republicans than for Democrats, supporting Hypothesis 1. For overall offensive speech, being a woman is associated with an approximately 4.2 percentage-point (ppt) increase in the predicted probability of offensive speech among white Republicans. By contrast, among white Democrats, the effect is negative, at roughly -2.8 ppt. This partisan contrast among white candidates is also visible across subcategories of offensive speech. For ridicule, white Republican women face an increase of about 3.4 ppt relative to white Republican men, while the corresponding effect among white Democrats is approximately -1.6 ppt. For demeaning speech, the effect is again positive among white Republicans, around 1.4 ppt, but negative among white Democrats, around -1.2 ppt. Hate speech follows a similar pattern: among white Republicans, being a woman is associated with about a 1.1 percentage-point increase in overall hate speech, compared with a near-zero or slightly positive effect among white Democrats.


Notably, among non-white candidates, the pattern is reversed, consistent with Hypothesis 2. For overall offensive speech, non-white Democratic women experience an increase of roughly 2.5 ppt relative to non-white Democratic men. By contrast, the estimate for non-white Republicans is negative, around -3 ppt, though the confidence interval is wide. A similar pattern appears for ridicule and demeaning speech. Among non-white Democrats, the gendered effect is positive, around 1.1 to 1.4 ppt, whereas among non-white Republicans, the estimates are negative or close to zero. For overall hate speech, the gendered effect among non-white Democrats is also strongly

positive, at approximately 1.2 ppt, while the estimate for non-white Republicans is much smaller and statistically uncertain.

The within-party comparisons also align with the hypotheses. Among Democrats, the gendered effect is stronger for non-white candidates than for white candidates, supporting Hypothesis 3. For overall offensive speech, the effect is about 2.5 ppt among non-white Democrats but approximately -2.8 ppt among white Democrats. For overall hate speech, the effect is about 1.2 ppt among non-white Democrats, compared with only about 0.05 ppt among white Democrats. The same pattern appears in identity-based hate speech categories: for race-based hate speech, the gendered effect is approximately 0.18 ppt among non-white Democrats, compared with near zero among white Democrats; for religion-based hate speech, it is about 0.11 ppt among non-white Democrats but close to zero among white Democrats.

Finally, among Republicans, the gendered effect is stronger for white candidates than for non-white candidates, supporting Hypothesis 4. White Republican women show positive gendered effects across most outcomes, including approximately 4.2 ppt for overall offensive speech, 3.4 points for ridicule, 1.4 points for demeaning speech, and 1.1 points for overall hate speech. In contrast, the corresponding estimates for non-white Republicans are generally smaller, negative, or imprecisely estimated. For example, the gendered effect for overall offensive speech among non-white Republicans is approximately -3 ppt, and for overall hate speech it is close to zero. This suggests that Republican women's heightened vulnerability is concentrated primarily among white candidates rather than non-white candidates.

Figure 3. Gendered Effect of Online Violence, by Candidate Party and Race




Note: The predicted probabilities of these results are derived from logistic regression models, with other control variables—such as Senate, Incumbency, Follower Count (log), Tweet Count (log), Impression count (log), and Time to Election—held at their mean values.

Figure 4 disaggregates the broader racial categories used in the previous analysis and shows that the patterns of online violence are not uniform within the non-white category. Among Democratic candidates, the highest predicted probabilities of offensive speech are observed for Asian women, at 21.8%, followed by Black women, at 17.6%. These rates are higher than those for their male co-partisans, especially among Asian Democrats, where the predicted probability for women is nearly 10 ppt higher than for men, 21.8% compared with 11.9%. The hate speech results show a similar but even more concentrated pattern. Black Democratic women have the highest predicted probability of hate speech in the figure, at 2.4%, compared with 0.3% for Black Democratic men. Asian Democratic women also face elevated hate speech, at 1.0%, compared with 0.4% for Asian Democratic men. By contrast, Hispanic Democratic women do not follow the same pattern: their predicted probability of offensive speech is 11.8%, lower than that of Hispanic Democratic men, while their predicted probability of hate speech is only modestly higher, 0.4% compared with 0.2%. This suggests that the online vulnerability of Democratic women of color is not evenly distributed across racial groups but is especially pronounced among Black and Asian women candidates.

The Republican side of the figure shows a different configuration of vulnerability. White Republican women stand out as the group most consistently exposed to both offensive and hateful content within the party: their predicted probability of offensive speech is 18.1%, compared with 14.1% for white Republican men, and their predicted probability of hate speech is 1.4%, compared with 0.1% for white Republican men. Among Republican candidates of color, however, the pattern is more uneven. Asian Republican women have a higher predicted probability of offensive speech than Asian Republican men, 15.7% compared with 12.4%, but the predicted probability of hate speech for Asian Republican women is close to zero. Hispanic Republican women show lower

predicted probabilities than Hispanic Republican men for both offensive speech, 11.7% compared with 20.9%, and hate speech, 0.2% compared with 0.5%. Black Republican women also show very low predicted probabilities, although this estimate should be interpreted cautiously given the small number of observations in that subgroup.

Figure 4. Predicted Probability of Offensive and Hate Speech by Gender, Race, and Party



Note: The predicted probabilities of these results are derived from logistic regression models, with other control variables—such as Senate, Incumbency, Follower Count (log), Tweet Count (log), Impression count (log), and Time to Election—held at their mean values.

These results sharpen the paper’s central argument by showing that the intersection of race, gender, and party produces distinct patterns of visibility and vulnerability rather than a single hierarchy of exposure. The figure suggests that Democratic Black and Asian women are especially likely to be targeted by identity-based hostility, while white Republican women are the most consistently vulnerable group within the Republican Party. At the same time, the relatively lower estimates for Hispanic women and several Republican women of color indicate that “women of color” should not be treated as a homogeneous category. These results therefore reinforce the need for an intersectional and partisan framework: online political violence is shaped not only by whether candidates are women or non-white, but by how specific racial identities acquire different meanings within Democratic and Republican partisan contexts.

Conclusion

This study demonstrates that online political violence in the contemporary U.S. is not structured solely by gender, race, or party affiliation in isolation, but rather by the intersection of all three. While women of color are often treated as a homogeneous group in public discourse, our analysis reveals that their exposure to digital abuse is highly contingent on partisan identity. Democratic women of color—particularly Asian and Black candidates—emerge as the most frequent targets of both offensive and hate speech, reflecting their symbolic representation of progressive, multicultural politics. These candidates, by virtue of their race, gender, and political affiliation, are perceived as challenging traditional hierarchies, making them high-salience targets for online hostility. In contrast, Republican women of color receive far less attention and are less frequently

attacked, likely because their political alignment does not position them as symbolic disruptors of racial and gender norms. Meanwhile, white Republican women experience a distinct form of gendered hostility, particularly when they are publicly aligned with patriarchal or anti-feminist rhetoric. These asymmetric patterns validate our four hypotheses and offer a nuanced understanding of how symbolic threat operates within polarized digital discourse.

Our findings further highlight that the intersectional nature of online violence disproportionately affects women compared to men, whose exposure to offensive and hate speech remains relatively low and consistent across race and party. This distinction underscores the unique vulnerabilities faced by female candidates, especially when their racial identity and partisan affiliation render them visible and contentious figures within political conflict. Unlike male candidates, who tend to face uniform patterns of online abuse regardless of race or party, women—particularly Democratic women of color and Republican white women—experience intensified hostility driven by their intersecting identities. Such disparities call for more nuanced analyses of political violence that move beyond simplistic categorizations of gender alone.

By integrating race, gender, and partisanship into the empirical study of online political violence, our research provides critical insights into how identity-based targeting operates in the digital public sphere. This intersectional approach challenges dominant frameworks that treat identity dimensions as additive rather than interactive, showing instead how political polarization shapes asymmetric vulnerabilities. Our findings carry important implications for democratic representation and political communication, as well as for social media platforms tasked with moderating harmful content. Recognizing how intersecting identities shape digital abuse can inform the development of more targeted and equitable content moderation practices. Moreover, by highlighting the unique risks faced by specific groups of women candidates, this study

underscores the broader societal challenge of ensuring equal political participation in an increasingly polarized and digitally mediated environment.

Bibliography

- Alzahrani, M. and AlGhamdi F. (2025). 'Social media sentiment analysis for sustainable rural event planning: A case study of agricultural festivals in Al-Baha, Saudi Arabia'. Sustainability, 17(9), 3864. <https://doi.org/10.3390/su17093864>.
- Amnesty International. (2018). Toxic Twitter: A Toxic Place for Women. Amnesty International, USA. Retrieved September 9, 2023 from <https://www.amnesty.org/en/latest/research/2018/03/online-violence-against-women-chapter-1/>
- Bailey, Moya. (2021). Misogynoir Transformed: Black Women's Digital Resistance. New York University Press, New York.
- Banet-Weiser, S., & Miltner, K. M. (2016). #MasculinitySoFragile: Culture, structure, and networked misogyny. *Feminist media studies*, 16(1), 171-174.
- Bardall, G. (2023). Policy Responses to Gender Based Political Violence Online. In *Gender and Violence Against Political Actors* (1st ed.), Bjarnegård, E., and Zetterberg, P. (eds.). Temple University Press, Philadelphia, PA, USA, 201–212.
- Bardall, G., Bjarnegård, E., & Piscopo, J. M. (2020). How is political violence gendered? Disentangling motives, forms, and impacts. *Political Studies*, 68(4), 916-935.
- Bejarano, C. E. 2013. *The Latino Gender Gap in U.S. Politics*. New York: Routledge.
- Bobo, L., and Gilliam, F. D. Jr. (1990). Race, Sociopolitical Participation, and Black Empowerment. *American Political Science Review*, 84 (2): 377–93.
- Bracic, A., Israel-Trummel, M., & Shortle, A. F. (2019). Is sexism for white people? Gender stereotypes, race, and the 2016 presidential election. *Political Behavior*, 41(2), 281-307.
- Brown, N. E., Clark, C. J., & Mahoney, A. M. (Eds.). (2023). *Women of Color Political Elites in the US: Intersectional Political Experiences*. Taylor & Francis.
- Camacho-Collados, J., Rezaee, K., Riahi, T., Ushio, A., Loureiro, D., Antypas, D., Boisson, J., Espinosa-Anke, L., Liu, F., Martínez-Cámara, E., Medina, G., Buhrmann, T., Neves, L., & Barbieri, F. (2022). TweetNLP: Cutting-Edge Natural Language Processing for Social Media. arXiv preprint arXiv:2206.14774.
- Carlson, C. R. (2019). Misogynistic hate speech and its chilling effect on women's free expression during the 2016 US Presidential campaign. *Journal of Hate Studies*, 14(1), 97-112.
- Cassese, E. C., & Barnes, T. D. (2019). Reconciling sexism and women's support for Republican candidates: A look at gender, class, and whiteness in the 2012 and 2016 presidential races. *Political Behavior*, 41, 677-700.

- Center for American Women and Politics (CAWP). (2024, September 12). Post-primary analysis: Women in 2024 congressional elections. Eagleton Institute of Politics, Rutgers University. <https://cawp.rutgers.edu/blog/post-primary-analysis-women-2024-congressional-elections>
- Collins, P. H. 1990. Black Feminist Thought: Knowledge, Consciousness, and the Politics of Empowerment. Hyman.
- Crenshaw, K. (1991). Mapping the Margins: Intersectionality, Identity Politics, and Violence Against Women of Color. *Stanford Law Review*, 43(6), 1241–1299.
- Dahlgren, P. (2005). The Internet, public spheres, and political communication: Dispersion and deliberation. *Political communication*, 22(2), 147-162.
- Erikson, J., Håkansson, S. and Josefsson, C. (2023). ‘Three dimensions of gendered online abuse: Analyzing Swedish MPs' experiences of social media.’ *Perspectives on Politics*, 21(3), 896–912. <https://doi.org/10.1017/S1537592721002048>.
- European Institute for Gender Equality. (2017). EIGE in brief 2017. <https://eige.europa.eu/publications-resources/publications/european-institute-gender-equality-eige-brief-2017>.
- Finkel, M., Thakur, D., Finkel, S., Duruşan, F., Yörük, E., Topçu I., Yardi, M. C., Zan, A., Han, J. (2025). Intersectional Hatred - An Application of Large Language Models to Detect Hate and Offensive Speech Targeted at Congressional Candidates in the 2024 U.S. Election. In Proceedings of WWW Companion '25. ACM, Sydney, NSW, Australia. 2776 – 2781.
- Evans, H. K. (2026). Issue Framing After Dobbs: Examining How Abortion Was Addressed in Tweets During the 2022 Midterm Elections. *PS: Political Science & Politics*, 1-9.
- Gershon, S. A., & Monforti, J. L. (2021). Intersecting campaigns: candidate race, ethnicity, gender and voter evaluations. *Politics, Groups, and Identities*.
- Gorrell, G., Bakir, M.E., Roberts, I., Greenwood, M.A. and Bontcheva, K. (2020). ‘Online abuse toward candidates during the UK General Election 2019.’ arXiv preprint, arXiv:2001.08686. <https://doi.org/10.48550/arXiv.2001.08686>.
- Hacıyakupoglu, G., & Wong, Y. (Eds.). (2022). *Gender and Security in Digital Space: Navigating Access, Harassment, and Disinformation*. Taylor & Francis.
- Han, J., Finkel, M., Finkel, S., Duruşan, F., and Yörük, E. (2026). “Online Gendered Violence and Visibility in Electoral Autocracies: Evidence from Parliamentary and Local Elections in Türkiye.” *European Journal of Political Research*, OnlineFirst.
- Harris-Perry, M. (2011). *Sister Citizen: Shame, Stereotypes, and Black Women in America*. New Haven, CT: Yale University Press.

- Hawkesworth, M. (2003). Congressional Enactments of Race–Gender: Toward a Theory of Raced–Gendered Institutions. *American Political Science Review*, 97(4), 529–550.
- Holm, M., Bjarnegård, E. and Zetterberg, P. (2024). ‘Comparing gendered exposure and impact in online election violence: Tunisian political candidates targeted on Facebook.’ *Politics & Gender*, 20(4), 879–902. <https://doi.org/10.1017/S1743923X24000163>.
- Jordan-Zachery, J. S. (2007). Am I a Black Woman or a Woman Who Is Black? A Few Thoughts on the Meaning of Intersectionality. *Politics & Gender*, 3(2), 254–263.
- Kwarteng, J., Farrell, T., Third, A., Fernandez, M. (2023) “Annotators' Perspectives: Exploring the Influence of Identity on Interpreting Misogynoir.” ASONAM '23: Proceedings of the International Conference on Advances in Social Networks Analysis and Mining. <https://doi.org/10.1145/3625007.3627292>
- Junn, J. (2017). The Trump majority: White womanhood and the making of female voters in the US. *Politics, Groups, and Identities*, 5(2), 343-352.
- Junn, J., & Masuoka, N. (2020). The gender gap is a race gap: Women voters in US presidential elections. *Perspectives on Politics*, 18(4), 1135-1145.
- Kim, C., & Junn, J. (2024). Whitewashing Women Voters: Intersectionality and Partisan Vote Choice in the 2020 US Presidential Election. *Politics & Gender*, 20(3), 701-726.
- Krook M. L. (2020). *Violence against Women in Politics*. Oxford University Press, Oxford, New York.
- Kuperberg, R. (2018). Intersectional violence against women in politics. *Politics & Gender*, 14(4), 685-690.
- Kuperberg, R. L., 2022. *World wide wrongs: Discrimination and online violence against women in politics*. PhD Thesis. Rutgers The State University of New Jersey, School of Graduate Studies, Rutgers.
- Lewis, R., Rowe, M., & Wiper, C. (2017). Online abuse of feminists as an emerging form of violence against women and girls. *The British Journal of Criminology*, 57(6), 1462–1481.
- Montoya, C. M., Bejarano, C., Brown, N. E., & Gershon, S. A. (2022). The intersectional dynamics of descriptive representation. *Politics & Gender*, 18(2), 483-512.
- Norwood, C., Jones, C., & Bolaji, L. (2021, June 17). More Black women are being elected to office. Few feel safe once they get there [PBS News Site]. PBS NewsHour. <https://www.pbs.org/newshour/politics/more-black-women-are-being-elected-to-office-few-feel-safe-once-they-get-there>

- Ortiz, R., Ahmed, S., Salas, P. and Eisty, N.U. (2025). ‘Analyzing social media engagement of computer science conferences.’ arXiv preprint, arXiv:2503.01038. <https://doi.org/10.48550/arXiv.2503.01038>.
- Pantoja, A. D., & Segura, G. M. (2003). Does ethnicity matter? Descriptive representation in legislatures and political alienation among Latinos. *Social Science Quarterly*, 84(2), 441-460.
- Pew Research Center. (2024, June 12). How Americans navigate politics on TikTok, X, Facebook and Instagram. <https://www.pewresearch.org/internet/2024/06/12/how-americans-navigate-politics-on-tiktok-x-facebook-and-instagram/>
- Phungula, N. (2023). Political violence directed towards female politicians in KZN Local Government, South Africa. *African Journal of Gender, Society & Development*, 12(3), 253.
- Rheault, L., Rayment, E., & Musulan, A. (2019). Politicians in the line of fire: Incivility and the treatment of women on social media. *Research & Politics*, 6(1), 2053168018816228.
- Schaffner, B. F. (2022). The heightened importance of racism and sexism in the 2018 US midterm elections. *British journal of political science*, 52(1), 492-500.
- Simien, E. M. (2007). Doing Intersectionality Research: From Conceptual Issues to Practical Examples. *Politics & Gender*, 3(2), 264–271.
- Smooth, Wendy G. (2006). Intersectionality in Electoral Politics: A Mess Worth Making. *Politics & Gender*, 2(3), 400–414.
- Sobieraj, S. (2020). *Credible threat: Attacks against women online and the future of democracy*. Oxford University Press.
- Thakur, D., Hankerson, D. L., Luria, M., Savage, S., Rodriguez, M., & Valdovinos, M. G. (2022). An Unrepresentative Democracy: How Disinformation and Online Abuse Hinder Women of Color Political Candidates in the United States (p. 95). Center for Democracy & Technology. <https://cdt.org/insights/an-unrepresentative-democracy-how-disinformation-and-online-abuse-hinder-women-of-color-political-candidates-in-the-united-states/>
- United Nations Human Rights Council. (2018). Report of the Working Group on the issue of discrimination against women in law and in practice: Note by the Secretariat (A/HRC/38/46). United Nations. <https://docs.un.org/en/A/HRC/38/46>
- United Nations. (2020). *Guidance on Addressing Hate Speech in the Field: Implementation of the United Nations Strategy and Plan of Action on Hate Speech*. New York: United Nations. Available at: https://digitallibrary.un.org/record/3889286/files/UN_Strategy_and_PoA_on_Hate_Speech_Guidance_on_Addressing_in_field.pdf, accessed 24 May 2024.

- Vidgen, B., Thrush, T., Waseem, Z., & Kiela, D. (2021). Learning from the Worst: Dynamically Generated Datasets to Improve Online Hate Detection. Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)
- Ward, L. M. (2016). Media and sexualization: State of empirical research, 1995–2015. *The Journal of Sex Research*, 53(4-5), 560-577.
- Wilfore, K. (2022). Security, Misogyny, and Disinformation Undermining Women's Leadership. In *Gender and Security in Digital Space* (pp. 124-142). Routledge.
- Zamfir, I. (2024). “Violence against women active in politics in the EU: A serious obstacle to political participation.” European Parliamentary Research Service, Briefing. [https://www.europarl.europa.eu/RegData/etudes/BRIE/2024/759600/EPRS_BRI\(2024\)759600_EN.pdf](https://www.europarl.europa.eu/RegData/etudes/BRIE/2024/759600/EPRS_BRI(2024)759600_EN.pdf)